

Intro

- We propose a generative two-level model that simultaneously represents 3D shapes using two levels of granularity, one for capturing fine-grained detail, the other for encoding a coarse structural decomposition.
- The two levels are tightly coupled via a shared latent space, wherein a single latent code vector decodes to two representations of the same shape. Modifications to one representation can be propagated to the other via the shared code.
- The shared latent space is learnt with a variational autoencoder (VAE). This approach
 1. imposes a Gaussian prior on the latent space, which enables sampling.
 2. encourages a compact latent space suitable for interpolation and optimization based manipulation.

Coarse Primitive-based shape representation,

- A SDF specifies, for every point $P = (P_x, P_y, P_z)$, the distance from that point to the nearest surface, where the sign encodes whether the point is inside or outside.
- Denote a set of N basic shape primitives by tuples:

$$\{(\mathbf{c}^i, \mathbf{q}^i) \mid i=1 \dots N\}$$

where \mathbf{c}^i describes the primitive type, and $\mathbf{q}^i \in \mathbb{R}^{k^i}$ describes the attributes of the primitives.

The dimensionality k^i denotes the dof for primitive i .

The SDF of a single element i is

$$d_{ci}(P, q^i) = SDF_{ci}(P, q^i)$$

An example of a simple geometric primitive is sphere;

$$k_{\text{sphere}} = 4$$

$$q^{\text{sphere}} = [c, r]$$

$c = (c_x, c_y, c_z)$ is center

r is radius

$$d_{\text{sphere}}(P, q^{\text{sphere}}) = \|P - c\|_2 - r$$

- to approximate the SDF of an arbitrarily complex shape, we construct the SDF of the union of geometric elements (spheres)

$$q = [q^1, \dots, q^N]$$

$$d_c(P, q) = \min_{1 \leq i \leq N} d_{ci}(P, q^i)$$

- To train the primitive-based model, given a target space X , e.g. mesh, sample pairs of 3D points P_t and their gt SDF $S_t = SDF_X(P_t)$, q can be learnt by minimizing the difference between predicted and real SDF:

$$\hat{q} = \arg \min_q \sum_t L_{SDF}(d_c(P_t, q), S_t).$$

High resolution shape representation.

Same with DeepSDF. directly learn the SDF with a neural network g_ϕ :

$$g_\phi(P) \approx SDF_X(P).$$

Learning a tightly coupled Latent Space.

- We learn a two-level shape representation over an entire class of shapes $\{X_j \mid j=1 \dots M\}$ by using two representation models that share the same latent code z_j .

- For representing multiple shapes with the primitive based coarse-level representation, we parameterize q with a neural network f_θ :

$$q_j = f_\theta(z_j)$$

f_θ is shared for all shapes.

For the fine-scale, we condition the neural network g_ϕ on latent code z_j :

$$g_\phi(z_j, P) \approx SDF_{X_j}(P).$$

VAD enforces a strong regularization on the latent space by representing the latent vector of each individual shape z_j with the parameters of its approximate posterior distribution (μ_j, σ_j) .

- We select the family of Gaussian distributions with diagonal covariance matrix as the approximate posterior of z , given shape X_j :

$$q(z \mid X=x_j) = \mathcal{N}(z; \mu_j, \sigma_j^2 \cdot I)$$

We apply the reparameterization trick, Sampling $\epsilon \sim \mathcal{N}(0, I)$ and setting $z_j = \mu_j + \sigma_j \odot \epsilon$ to allow optimization of μ_j, σ_j via GD.

During training, we max. the lower bound of the marginal likelihood (ELBO) over the dataset, which is the sum over lower bound of each individual shape X :

$$\log P_{\theta, \phi}(x) \geq \mathbb{E}_{z \sim q(z \mid X)} [\log P_{\theta, \phi}(x \mid z)] - D_{KL}(q(z \mid X) \parallel P(z)).$$

The learnable parameters are θ, ϕ , and variational parameters $\{(\mu_j, \sigma_j) \mid j=1 \dots M\}$, that parameterize $q(z \mid X)$.

- Since we'd like the two representations to be tightly coupled, i.e. both assign high probability density to a shape X_j given its latent code $z_j \sim q(z \mid X=x_j)$, we use a mixture model:

$$P_{\theta, \phi}(x \mid z) = \frac{P_\theta(x \mid z) + P_\phi(x \mid z)}{2}$$

$P_\theta(x \mid z), P_\phi(x \mid z)$ are the posterior distributions of coarse and fine representations.

$$\log P_\theta(x \mid z) = -\lambda_1 \int p(p) L_{SDF}(d_c(p, f_\theta(z)), SDF_X(p)) dp \quad \}$$

$$\log P_\phi(x \mid z) = -\lambda_2 \int p(p) L_{SDF}(g_\phi(z, p), SDF_X(p)) dp. \quad \}$$

Eq. above can be approximated via Monte Carlo, where p is sampled randomly from the 3D space following a specific rule $p(p)$.